


6.036 Project 2: Handwritten Digit Recognition

1) Multinomial/Softmax Regression and Gradient Descent

Final Test Error = 0.1005.


Graph of Cost Function Progression:




2) Using Manually Crafted Features

2. Test Error = 0.136.

Graph of Cost Function Progression:




3. Visualization of the first 100 MNIST images, as represented in the space by the first 2 principal components of the training data:



4. Plots of the reconstructions of the first two MNIST images (from their 20-dimensional PCA-representations) alongside the originals:






5. $x = [x_1, x_2]$ so that

$$\begin{aligned} (1 + x \cdot x)^2 &= (1 + x_1 x'_1 + x_2 x'_2)^2 \\ &= 1 + 2(x_1 x'_1 + x_2 x'_2) + (x_1 x'_1)^2 + 2(x_1 x'_1)(x_2 x'_2) + (x_2 x'_2)^2 \\ &= [1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, \sqrt{2}x_1 x_2, x_2^2] \cdot [1, \sqrt{2}x'_1, \sqrt{2}x'_2, x_1'^2, \sqrt{2}x'_1 x'_2, x_2'^2] \end{aligned}$$

So, $\phi(x) = [1, \sqrt{2}x_1, \sqrt{2}x_2, x_1^2, \sqrt{2}x_1 x_2, x_2^2]^T$. The `quadraticFeatures` function is implemented in `features.py`.

6. Test Error = 0.0368.

Graph of Cost Function Progression:



3) Classification Using Deep Neural Networks

1. (b) I got a test accuracy of 98.29%. I increased the number of neurons in the first fully-connected layer to 512, and I increased the learning rate to 0.1. I realized that increasing the learning rate by itself (without changing any other parameters) wasn't enough to get an accuracy above 98%. However, when tweaked together with the number of neurons in the first hidden layer, I observed a better test accuracy.

2. (a) The code for my CNN is in `mnist_nnet_cnn_skeleton.py`. It achieves a training accuracy of 94.50% and a validation accuracy of 98.36%.